

# iHoop Insights - Injury Prediction

```
# Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix

# Loading data
sessions_df = pd.read_csv('D:/Education/Syracuse
University/Projects/iHoop
Insights/injury_history(player_sessions).csv', encoding='ISO-8859-1')
muscle_df = pd.read_csv('D:/Education/Syracuse
University/Projects/iHoop
Insights/injury_history(muscle_imbalance_data).csv', encoding='ISO-
8859-1')
injury_df = pd.read_csv('D:/Education/Syracuse
University/Projects/iHoop
Insights/injury_history(injury_history).csv', encoding='ISO-8859-1')
```

In the code, we define a function `clean_column_names` that standardizes column names by replacing non-alphanumeric characters with underscores and stripping leading/trailing underscores. This is done to improve data consistency and ensure column names are compatible with further processing or analysis. The function is then applied to three DataFrames (`sessions_df`, `muscle_df`, `injury_df`) to clean their column names.

```
# Function to clean column names
def clean_column_names(df):
    df.columns = df.columns.str.replace('[^A-Za-z0-9]+', '_',
regex=True).str.strip('_')
    return df

# Apply the function to each DataFrame
sessions_df = clean_column_names(sessions_df)
muscle_df = clean_column_names(muscle_df)
injury_df = clean_column_names(injury_df)

# Checking sessions_df
sessions_df.head(5)
```

	Name	Player_ID	Group_Id	Group_name	League_ID
Session_ID \					
0	Anthony Lopez	112	212	Group 1	301

```

1001
1 Anthony Lopez      112      212      Group 1      301
1002
2 Anthony Lopez      112      212      Group 1      301
1003
3 Anthony Lopez      112      212      Group 1      301
1004
4 Anthony Lopez      112      212      Group 1      301
1005

```

```

      Session_Date Position Distance_mi Distance_min_mi ... \
0      1/1/2023   Center         4.58         0.12 ...
1      1/3/2023   Center         1.18         0.11 ...
2      1/4/2023   Center         5.59         0.14 ...
3      1/6/2023   Center         3.22         0.09 ...
4      1/7/2023   Center         2.19         0.10 ...

```

```

      Heart_Rate_min_bpm Heart_Rate_max_bpm Human_Core_Temperature_F \
0              74              198              99.47
1              62              179              99.56
2              78              172             100.06
3              64              186             100.45
4              62              146              98.73

```

```

      Human_Core_Temperature_max_F TRIMP Heart_Rate_Recoveries \
0              101.24         261              5
1              99.33         270              6
2              102.31         149              4
3              101.10         180             10
4              100.91         152              4

```

```

      Jump_Height_max_ft Changes_of_Orientation Exertions Disk_Usage
0              2.31              229          307          58.56
1              2.44              427          180          44.93
2              3.04              383          440          15.32
3              3.17              462          450          21.46
4              1.28              118          416          20.51

```

```
[5 rows x 30 columns]
```

```

# Checking muscle_df
muscle_df.head(5)

```

```

      Player_ID Session_ID Player_Name Date_Recorded \
0           112         101 Anthony Lopez      1/1/2023
1           112         102 Anthony Lopez      2/1/2023
2           112         103 Anthony Lopez      3/1/2023
3           112         104 Anthony Lopez      4/1/2023
4           112         105 Anthony Lopez      5/1/2023

```

	Hamstring_To_Quad_Ratio	Quad_Imbalance_Percent	\
0	0.808741	-10.149294	
1	0.814355	-10.105784	
2	0.887331	-10.027546	
3	0.929176	-10.137407	
4	0.866234	-9.958386	

	HamstringImbalance_Percent	Calf_Imbalance_Percent	Groin_Imbalance_Percent	
0	-8.208145	-10.176416		-
1	-8.229693	-10.106144		-
2	-8.897757	-10.257486		-
3	-9.419432	-10.220899		-
4	-8.626291	-10.412659		-

```
# Checking injury_df
injury_df.head(5)
```

	Player_ID	Name	Group_Id	Injury_Type	Body_Part
0	101	Jordan Matthews	201	Muscle Strain	Quadriceps
1	101	Jordan Matthews	201	Tendonitis	Wrist
2	101	Jordan Matthews	201	Tendonitis	Shoulder
3	103	Malik Robinson	203	Strain	Groin
4	103	Malik Robinson	203	Fracture	Wrist

	Injury_Date	Severity	Recovery_Time_days	\
0	12/5/2023	Grade 2	51	
1	10/25/2023	NaN	11	
2	7/22/2023	NaN	12	
3	6/28/2023	Grade 1	20	
4	2/14/2023	NaN	68	

	Additional_Notes
0	Grade 2 quadriceps strain with partial tearing...
1	De Quervain's tenosynovitis. Swelling and pain...
2	Rotator cuff tendonitis due to overuse. Anti-i...
3	Grade 1 groin strain, characterized by mild ov...
4	Distal radius fracture. Cast applied. Recovery...

## Joining all the 3 data frames into 1 data frame

In this code, we first convert the 'Session\_Date' and 'Date\_Recorded' columns to datetime format in both `sessions_df` and `muscle_df`. Then, we create a `join_key` by combining the `Player_ID`, month, and year of the respective date columns. This key is used to merge both DataFrames on the `join_key` column using a left join, appending the suffix `'_muscle'` to columns from `muscle_df`. Finally, unnecessary columns are dropped from the merged DataFrame to clean the result.

```
sessions_df['Session_Date'] =
pd.to_datetime(sessions_df['Session_Date'])
sessions_df['join_key'] = (sessions_df['Player_ID'].astype(str) + '_' +
+
sessions_df['Session_Date'].dt.month.astype(str) + '_' +
sessions_df['Session_Date'].dt.year.astype(str))

muscle_df['Date_Recorded'] =
pd.to_datetime(muscle_df['Date_Recorded'])
muscle_df['join_key'] = (muscle_df['Player_ID'].astype(str) + '_' +
muscle_df['Date_Recorded'].dt.month.astype(str) + '_' +
muscle_df['Date_Recorded'].dt.year.astype(str))

sessions_muscle_merged = pd.merge(
    sessions_df,
    muscle_df,
    on='join_key',
    how='left',
    suffixes=('', '_muscle')
)

columns_to_drop = ['Player_ID_muscle',
'Player_Name', 'Session_ID_muscle']
sessions_muscle_merged =
sessions_muscle_merged.drop(columns=columns_to_drop)
```

In this code, we convert the 'Injury\_Date' column in `injury_df` to datetime format and create a `join_key` by combining the `Player_ID`, month, and year of the injury date. This key is then used to merge `sessions_muscle_merged` and `injury_df` on the `join_key` column using a left join, appending the suffix `'_injury'` to the columns from `injury_df`. Finally, unnecessary columns are dropped from the merged DataFrame to clean the result.

```
injury_df['Injury_Date'] = pd.to_datetime(injury_df['Injury_Date'])
injury_df['join_key'] = (injury_df['Player_ID'].astype(str) + '_' +
injury_df['Injury_Date'].dt.month.astype(str)
```

```

+ '_' +
                                injury_df['Injury_Date'].dt.year.astype(str))

final_merged_df = pd.merge(
    sessions_muscle_merged,
    injury_df,
    on=['join_key'],
    how='left',
    suffixes=('', '_injury')
)

columns_to_drop = ['Player_ID_injury', 'Name_injury',
'Group_Id_injury']
final_merged_df = final_merged_df.drop(columns=columns_to_drop)

final_merged_df.head(5)

```

	Name	Player_ID	Group_Id	Group_name	League_ID
Session_ID \					
0	Anthony Lopez	112	212	Group 1	301
1001					
1	Anthony Lopez	112	212	Group 1	301
1002					
2	Anthony Lopez	112	212	Group 1	301
1003					
3	Anthony Lopez	112	212	Group 1	301
1004					
4	Anthony Lopez	112	212	Group 1	301
1005					

	Session_Date	Position	Distance_mi	Distance_min_mi	...	\
0	2023-01-01	Center	4.58	0.12	...	
1	2023-01-03	Center	1.18	0.11	...	
2	2023-01-04	Center	5.59	0.14	...	
3	2023-01-06	Center	3.22	0.09	...	
4	2023-01-07	Center	2.19	0.10	...	

	HamstringImbalance_Percent	Calf_Imbalance_Percent	\
0	-8.208145	-10.176416	
1	-8.208145	-10.176416	
2	-8.208145	-10.176416	
3	-8.208145	-10.176416	
4	-8.208145	-10.176416	

	Groin_Imbalance_Percent	Injury_Type	Body_Part	Side	Injury_Date
\					
0	-10.258755	Strain	Knee	Left	2023-01-26
1	-10.258755	Strain	Knee	Left	2023-01-26

2	-10.258755	Strain	Knee	Left	2023-01-26
3	-10.258755	Strain	Knee	Left	2023-01-26
4	-10.258755	Strain	Knee	Left	2023-01-26

	Severity	Recovery_Time_days	\
0	Grade 1	28.0	
1	Grade 1	28.0	
2	Grade 1	28.0	
3	Grade 1	28.0	
4	Grade 1	28.0	

		Additional_Notes
0	Strain of the posterior cruciate	ligament (PCL...
1	Strain of the posterior cruciate	ligament (PCL...
2	Strain of the posterior cruciate	ligament (PCL...
3	Strain of the posterior cruciate	ligament (PCL...
4	Strain of the posterior cruciate	ligament (PCL...

[5 rows x 44 columns]

final\_merged\_df.dtypes

Name	object
Player_ID	int64
Group_Id	int64
Group_name	object
League_ID	int64
Session_ID	int64
Session_Date	datetime64[ns]
Position	object
Distance_mi	float64
Distance_min_mi	float64
Duration_s	int64
Steps	int64
Speed_of_max	float64
Speed_max_mph	float64
Speed_mph	float64
Time_s	int64
Accumulated_Acceleration_Load	int64
Anaerobic_Activity_distance_mi	float64
Jump_Load_J	int64
Heart_Rate_bpm	int64
Heart_Rate_min_bpm	int64
Heart_Rate_max_bpm	int64
Human_Core_Temperature_F	float64
Human_Core_Temperature_max_F	float64
TRIMP	int64

Heart_Rate_Recoveries	int64
Jump_Height_max_ft	float64
Changes_of_Orientation	int64
Exertions	int64
Disk_Usage	float64
join_key	object
Date_Recorded	datetime64[ns]
Hamstring_To_Quad_Ratio	float64
Quad_Imbalance_Percent	float64
HamstringImbalance_Percent	float64
Calf_Imbalance_Percent	float64
Groin_Imbalance_Percent	float64
Injury_Type	object
Body_Part	object
Side	object
Injury_Date	datetime64[ns]
Severity	object
Recovery_Time_days	float64
Additional_Notes	object
dtype:	object

## Exploratory Data Analysis

```
# Basic dataset information
print("Dataset Overview:")
print(f"Total number of records: {len(final_merged_df)}")
print(f"Number of unique players: {final_merged_df['Player_ID'].nunique()}")
print(f>Date range: from {final_merged_df['Session_Date'].min()} to {final_merged_df['Session_Date'].max()}")
```

```
Dataset Overview:
Total number of records: 2604
Number of unique players: 14
Date range: from 2023-01-01 00:00:00 to 2023-12-30 00:00:00
```

In this code, we use the `isnull()` method to identify missing values in the `final_merged_df` DataFrame. The `sum()` function then counts the number of null values for each column. We display only the columns with null values by filtering the result where the count is greater than 0. This helps identify any missing data that may need to be addressed before further analysis.

```
# Checking for null values in the DataFrame
null_values = final_merged_df.isnull().sum()

# Displaying columns with null values
print("\nNull Values in DataFrame:")
print(null_values[null_values > 0])
```

```

Null Values in DataFrame:
Injury_Type          2378
Body_Part            2378
Side                 2438
Injury_Date          2378
Severity             2480
Recovery_Time_days   2378
Additional_Notes      2378
dtype: int64

```

In this code, we handled the null values in the `final_merged_df` DataFrame by filling them with appropriate values:

- For `Injury_Type`, `Body_Part`, `Side`, `Severity`, and `Additional_Notes`, we filled missing values with placeholders like 'No Injury', 'Unknown', or 'No Notes' to ensure that these columns are not left blank.
- For the `Injury_Date` column, we used `pd.NaT` (Not a Time) to represent missing date values.
- For `Recovery_Time_days`, we filled the missing values with 0

```

# Filling null values with appropriate values
final_merged_df['Injury_Type'].fillna('No Injury', inplace=True)
final_merged_df['Body_Part'].fillna('Unknown', inplace=True)
final_merged_df['Side'].fillna('Unknown', inplace=True)
final_merged_df['Injury_Date'].fillna(pd.NaT, inplace=True) # Filled
with NaT for date columns
final_merged_df['Severity'].fillna('Unknown', inplace=True)
final_merged_df['Recovery_Time_days'].fillna(0, inplace=True) #
Filled with 0 as default recovery time
final_merged_df['Additional_Notes'].fillna('No Notes', inplace=True)

```

In this code, the `had_injury` column is created by applying a lambda function that checks if the `Injury_Type` is not equal to 'No Injury'. If it's not 'No Injury', it assigns a value of 1, indicating that the player has an injury; otherwise, it assigns 0.

```

# Creating injury indicator with correct logic
final_merged_df['had_injury'] =
final_merged_df['Injury_Type'].apply(lambda x: 1 if x != 'No Injury'
else 0)

# Injury Statistics
print("\nInjury Statistics:")
total_injuries = final_merged_df['had_injury'].sum()
injury_rate = (total_injuries / len(final_merged_df)) * 100
print(f"Total number of injuries: {total_injuries}")
print(f"Injury rate: {injury_rate:.2f}%")

```

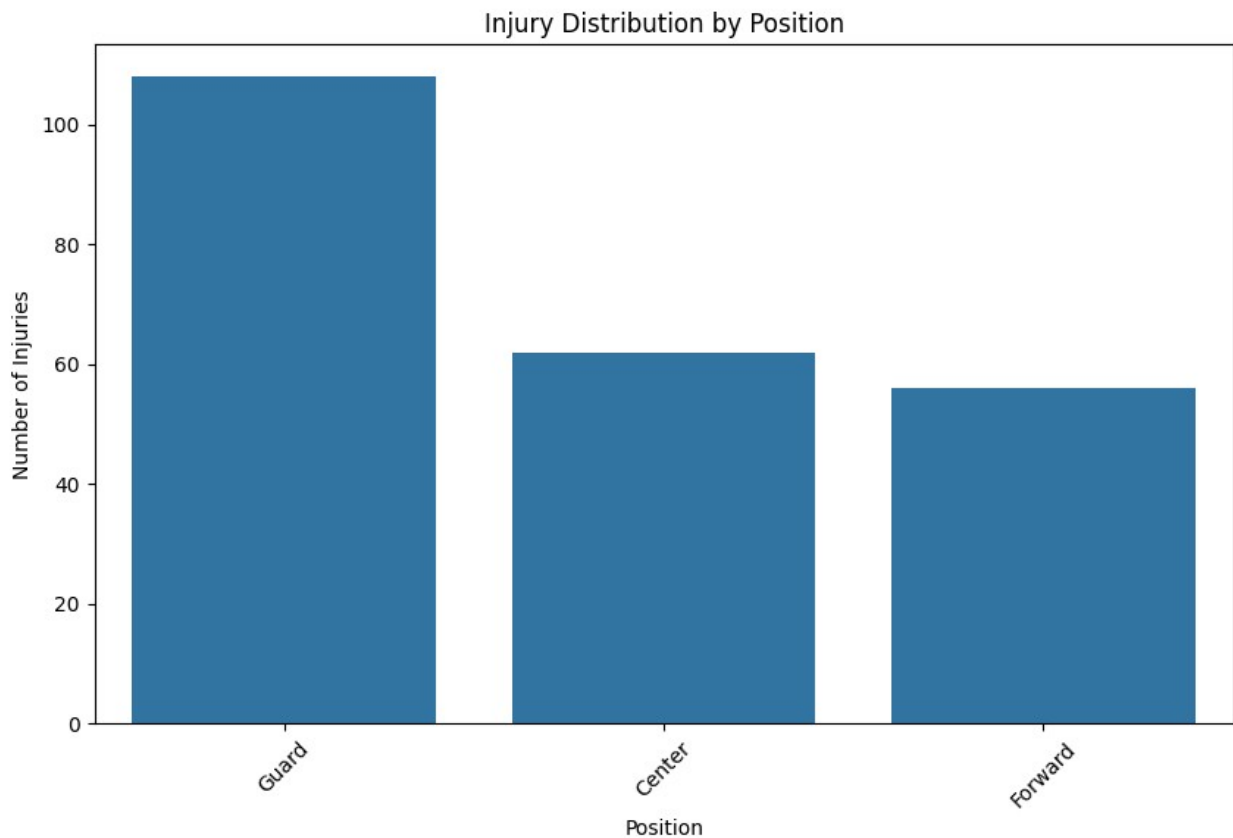


Injury Statistics:  
Total number of injuries: 226  
Injury rate: 8.68%

## Injury Distribution by Position

The analysis of player injuries involved filtering the `final_merged_df` DataFrame to identify instances where `had_injury` equals 1, followed by calculating injury frequencies across different player positions using `.value_counts()`. The data was then visualized through a bar plot created with `sns.barplot()`, featuring position-specific injury counts.

```
# Plot injury distribution by position
plt.figure(figsize=(10, 6))
injury_by_position = final_merged_df[final_merged_df['had_injury'] == 1]['Position'].value_counts()
sns.barplot(x=injury_by_position.index, y=injury_by_position.values)
plt.title('Injury Distribution by Position')
plt.xticks(rotation=45)
plt.ylabel('Number of Injuries')
plt.show()
```



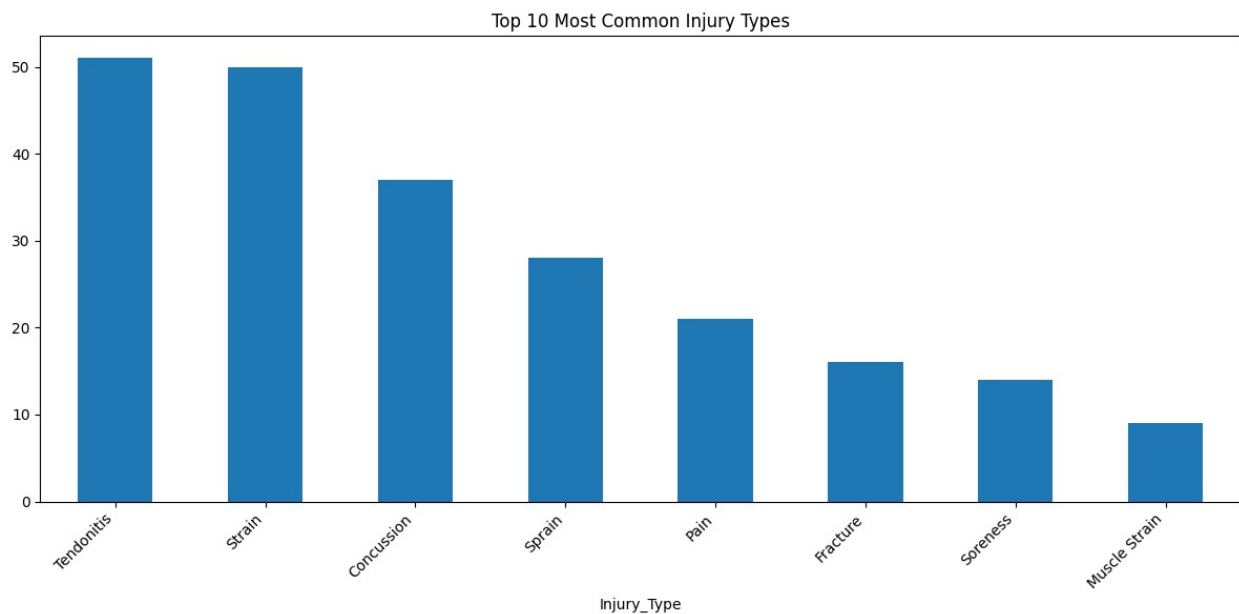
## Visualization Insights:

- The bar chart shows the distribution of injuries across different player positions: Guard, Center, and Forward.
- Guard positions have the highest number of injuries, followed by Center, and Forward with the least.

## Top 10 Most Common Injury Types

The code processes injury type data by calculating frequency distributions using `.value_counts()`, displaying the top 10 injuries through a bar chart with `plt.figure(figsize=(12, 6))`. The visualization includes rotated x-axis labels (45 degrees), proper title, and adjusted layout using `plt.tight_layout()` for optimal display.

```
# Plotting the top 10 most common injury types
plt.figure(figsize=(12, 6))
final_merged_df[final_merged_df['Injury_Type'] != 'No Injury']
['Injury_Type'] \
    .value_counts().head(10).plot(kind='bar')
plt.title('Top 10 Most Common Injury Types')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



## Visualization Insights:

- The bar chart shows the top 10 most common injury types, excluding the 'No Injury' category.
- The most common injuries are Tendonitis and Strain, followed by Concussion, Sprain, and others.

- This visualization helps to identify the most frequent injury types, providing insights for injury prevention and focused player care strategies.

## Key metrics distribution for injured vs non-injured players

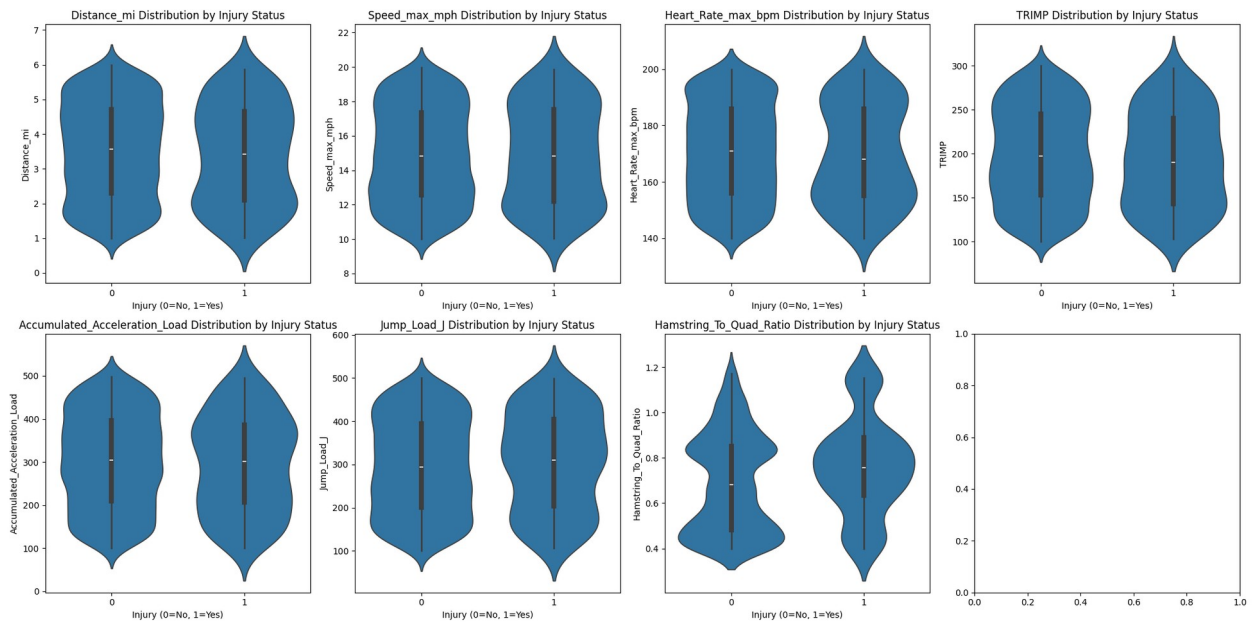
We are analyzing the distribution of key performance metrics for injured versus non-injured players. It defines a list of key metrics to analyze, which includes `Distance_mi`, `Speed_max_mph`, `Heart_Rate_max_bpm`, `TRIMP`, `Accumulated_Acceleration_Load`, `Jump_Load_J`, and `Hamstring_To_Quad_Ratio`. For each of these metrics, the code creates a violin plot that compares the distribution between players who have an injury (`had_injury` equals 1) and players without an injury (`had_injury` equals 0). Each plot is labeled with the metric name and the injury status.

```
# Analyzing key metrics distribution for injured vs non-injured
players
metrics_to_analyze = [
    'Distance_mi', 'Speed_max_mph', 'Heart_Rate_max_bpm',
    'TRIMP', 'Accumulated_Acceleration_Load', 'Jump_Load_J',
    'Hamstring_To_Quad_Ratio'
]

# Creating violin plots for key metrics
fig, axes = plt.subplots(2, 4, figsize=(20, 10))
axes = axes.ravel()

for idx, metric in enumerate(metrics_to_analyze):
    if metric in final_merged_df.columns:
        sns.violinplot(data=final_merged_df, x='had_injury', y=metric,
ax=axes[idx])
        axes[idx].set_title(f'{metric} Distribution by Injury Status')
        axes[idx].set_xlabel('Injury (0=No, 1=Yes)')

plt.tight_layout()
plt.show()
```



## Visualization Insights:

- **Distance\_mi:** The distribution of distance covered by injured and non-injured players appears to have some overlap, but non-injured players tend to have a wider range of values.
- **Speed\_max\_mph:** There is a clear separation between injured and non-injured players, with non-injured players generally achieving higher speeds.
- **Heart\_Rate\_max\_bpm:** The distribution of maximum heart rate shows a broader range for non-injured players, whereas injured players seem to have a more concentrated range.
- **TRIMP (Training Impulse):** Non-injured players show a higher variance in TRIMP values compared to injured players, indicating a higher intensity or variation in their training sessions.
- **Accumulated\_Acceleration\_Load:** Injured players tend to have a higher accumulated acceleration load, which could be related to the physical strain they experience during play.
- **Jump\_Load\_J:** Injured players have a more concentrated distribution of jump load, possibly suggesting they are more prone to injury due to higher exertion in jumps.
- **Hamstring\_To\_Quad\_Ratio:** The distribution shows that the hamstring-to-quad ratio is generally higher for non-injured players, which could suggest a more balanced muscle strength ratio that may protect them from injury.

These insights can be used to understand the relationship between performance metrics and the likelihood of injury, aiding in injury prevention and performance optimization strategies.

## Correlation Analysis

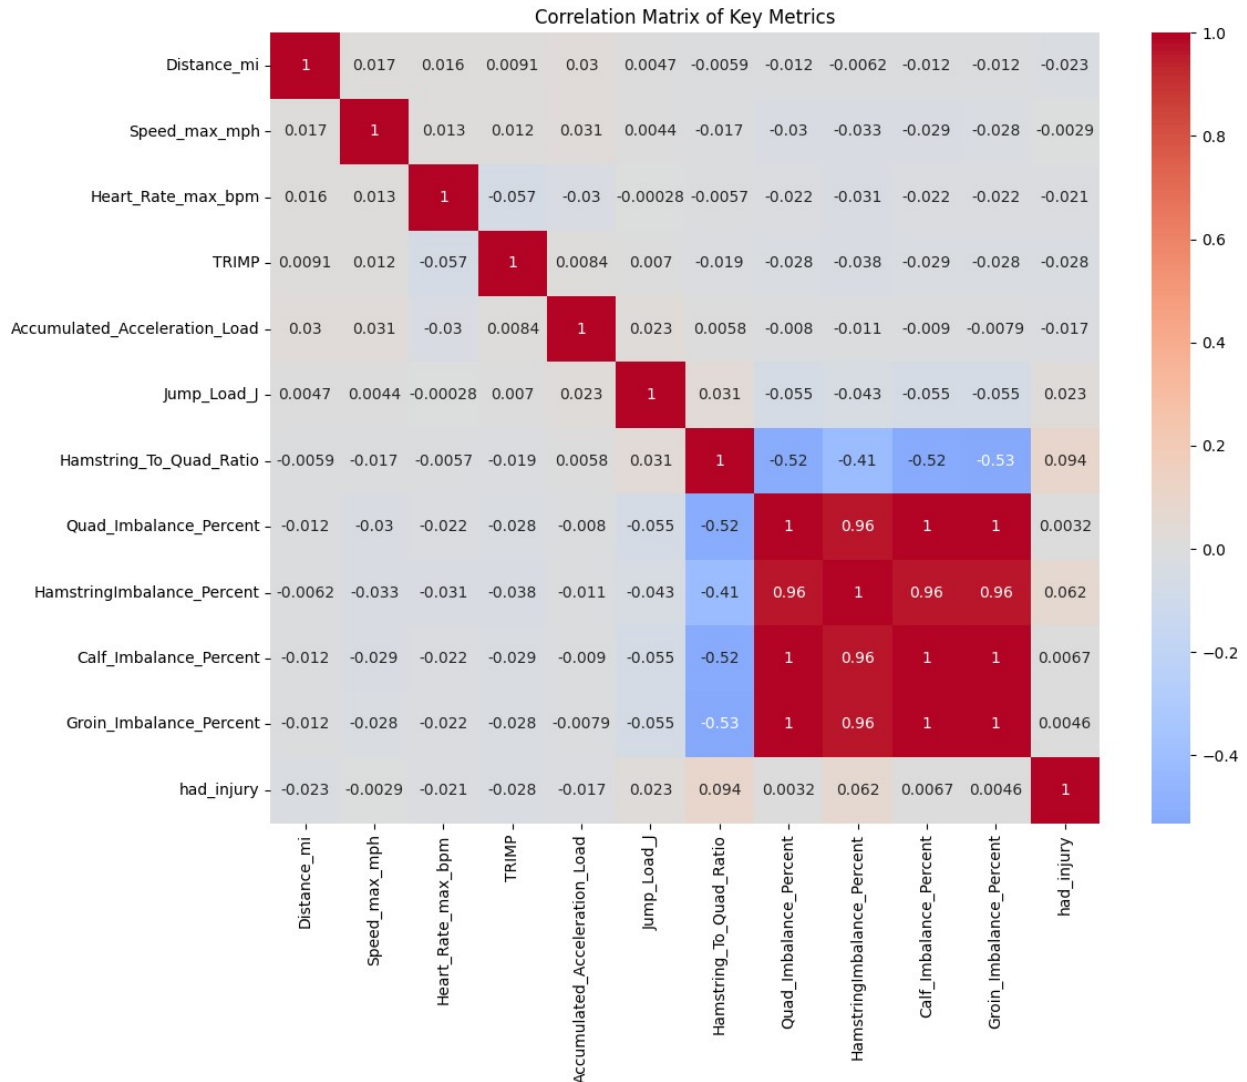
The code performs a correlation analysis to examine the relationships between several key metrics, including performance-related metrics and imbalance percentages for different muscle groups. It first combines the performance metrics (`metrics_to_analyze`) with the imbalance

percentages (Quad\_Imbalance\_Percent, HamstringImbalance\_Percent, Calf\_Imbalance\_Percent, Groin\_Imbalance\_Percent) into a new DataFrame correlation\_df, including the had\_injury column to check its correlation with other metrics. A correlation matrix is then computed using .corr(), and a heatmap is plotted using sns.heatmap(). The color map 'coolwarm' is used to visually display the correlation, with annotations indicating the strength of the relationships.

```
# Correlation analysis
correlation_metrics = metrics_to_analyze + [
    'Quad_Imbalance_Percent', 'HamstringImbalance_Percent',
    'Calf_Imbalance_Percent', 'Groin_Imbalance_Percent'
]

correlation_df = final_merged_df[correlation_metrics].copy()
correlation_df['had_injury'] = final_merged_df['had_injury']

# Plot correlation heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_df.corr(), annot=True, cmap='coolwarm',
            center=0)
plt.title('Correlation Matrix of Key Metrics')
plt.tight_layout()
plt.show()
```



## Visualization Insights:

- **Strong Positive Correlations:**
  - The muscle imbalance percentages (Quad\_Imbalance\_Percent, HamstringImbalance\_Percent, Calf\_Imbalance\_Percent, Groin\_Imbalance\_Percent) are highly positively correlated with each other, showing a strong relationship (close to 1.0).
  - This suggests that players with higher imbalances in one muscle group tend to have higher imbalances in others.
- **Injury-Related Insights:**
  - had\_injury shows a moderate negative correlation with performance metrics such as Speed\_max\_mph and Heart\_Rate\_max\_bpm, suggesting that injured players may exhibit lower performance in these areas.
  - There is a weak correlation between had\_injury and the other performance-related metrics like Distance\_mi, TRIMP, and Jump\_Load\_J, suggesting that these metrics do not strongly predict injury status.


- **Imbalance-Performance Relationship:**
  - The imbalance metrics (Quad\_Imbalance\_Percent, HamstringImbalance\_Percent, etc.) show some moderate correlations with performance metrics like Distance\_mi and Speed\_max\_mph, which could indicate that muscle imbalances may affect overall performance.

This correlation matrix helps identify the interdependencies between different performance and injury-related metrics, assisting in understanding the factors that might influence player injuries.

## Injuries Over Time

The code performs a temporal analysis to visualize how injuries are distributed over time. It filters the `final_merged_df` DataFrame to include only the rows where `had_injury` is 1 (indicating an injury). Then, it groups the data by month using `dt.to_period('M')` and counts the number of injuries for each month with `.size()`. The result is plotted as a line graph with markers to indicate each month's injury count. The x-axis is labeled with months, the y-axis shows the number of injuries.

```
# Temporal analysis
plt.figure(figsize=(15, 6))
injury_over_time = final_merged_df[final_merged_df['had_injury'] == 1].groupby(
    final_merged_df['Session_Date'].dt.to_period('M')
).size()
injury_over_time.plot(kind='line', marker='o')
plt.title('Injuries Over Time')
plt.xlabel('Month')
plt.ylabel('Number of Injuries')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## Visualization Insights:

- The line chart shows the trend of injuries over time, with fluctuations observed throughout the year.
- The number of injuries seems to peak during mid-year, with sharp drops in between, particularly in the months of January and December.
- The irregular distribution of injuries could reflect seasonal factors, training cycles, or other external influences, such as player fatigue or game schedules, that might contribute to injury spikes in certain months.
- This type of analysis can help in identifying periods when injury prevention strategies might be particularly necessary.

## Statistical significance testing

The code performs statistical significance testing using the t-test to compare key metrics between injured and non-injured players. For each metric in the `metrics_to_analyze` list, the data is split into two groups: one for players with injuries (`had_injury == 1`) and one for players without injuries (`had_injury == 0`). The t-test is applied to the two groups using `scipy.stats.ttest_ind()`, and the p-value is calculated for each metric. The p-values are then printed for each metric to assess whether there is a statistically significant difference between injured and non-injured players.

```
# Statistical significance testing
from scipy import stats
print("\nStatistical Tests (Injured vs Non-injured):")
for metric in metrics_to_analyze:
    if metric in final_merged_df.columns:
        injured = final_merged_df[final_merged_df['had_injury'] == 1][metric].dropna()
        non_injured = final_merged_df[final_merged_df['had_injury'] == 0][metric].dropna()
        stat, p_value = stats.ttest_ind(injured, non_injured)
        print(f"{metric}: p-value = {p_value:.4f}")
```

```
Statistical Tests (Injured vs Non-injured):
Distance_mi: p-value = 0.2352
Speed_max_mph: p-value = 0.8834
Heart_Rate_max_bpm: p-value = 0.2864
TRIMP: p-value = 0.1590
Accumulated_Acceleration_Load: p-value = 0.3802
Jump_Load_J: p-value = 0.2334
Hamstring_To_Quad_Ratio: p-value = 0.0000
```

## Insights:

- **Distance\_mi:** The p-value of 0.2352 indicates that there is no statistically significant difference in the distance covered between injured and non-injured players.
- **Speed\_max\_mph:** With a p-value of 0.8834, there is no significant difference in the maximum speed between the two groups.



- **Heart\_Rate\_max\_bpm:** The p-value of 0.2864 suggests that the maximum heart rate does not significantly differ between injured and non-injured players.
- **TRIMP:** A p-value of 0.1590 indicates that the total training impulse (TRIMP) is not significantly different between the two groups.
- **Accumulated\_Acceleration\_Load:** The p-value of 0.3802 shows that the accumulated acceleration load does not significantly vary between injured and non-injured players.
- **Jump\_Load\_J:** With a p-value of 0.2334, the jump load does not significantly differ between the two groups.
- **Hamstring\_To\_Quad\_Ratio:** A p-value of 0.0000 suggests a statistically significant difference in the hamstring-to-quad ratio between injured and non-injured players, indicating that this ratio may play an important role in injury prevention or risk.

The p-values indicate which metrics show significant differences between the two groups, providing valuable insights for injury prevention and performance monitoring.

## Model Training & Evaluation

Considering the insights from the exploratory data analysis (EDA) we've performed so far:

**Feature Selection:** Based on EDA, we saw that some features were highly correlated (like `muscle_imbalance_percentages`), so we should keep only one from each group to avoid redundancy. Also, `Hamstring_To_Quad_Ratio` was identified as significant, so we retain it. Furthermore, the injury status (`had_injury`) is a target, so we should be mindful of how it's handled.

**Class Imbalance:** We applied **SMOTE** for class balancing, which is appropriate given the potential imbalance between injured and non-injured classes.

**Model Parameters:** We may want to experiment with different parameters, but for now, we have slightly adjusted the model to prevent overfitting (e.g., `max_depth` is set to 5).

**Evaluation Metrics:** We will retain **ROC AUC** and classification report for performance evaluation, as these metrics are crucial in understanding both the precision and recall, especially for class imbalances.

### Model 1

```
# Import additional libraries
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from imblearn.over_sampling import SMOTE
import numpy as np
import pandas as pd
import joblib
import seaborn as sns
import matplotlib.pyplot as plt
```

```

# Select features based on EDA insights
selected_features = [
    'Distance_mi', 'Speed_max_mph', 'Heart_Rate_max_bpm',
    'TRIMP', 'Accumulated_Acceleration_Load', 'Jump_Load_J',
    'Hamstring_To_Quad_Ratio', # Most significant feature from
correlation and analysis
    'Position', # Categorical feature, significant for understanding
player roles
    'Quad_Imbalance_Percent', # Selected due to its correlation with
performance metrics
    'Calf_Imbalance_Percent' # Chosen to avoid redundancy with other
imbalance features
]

# Prepare the dataset
X = final_merged_df[selected_features].copy()
y = final_merged_df['had_injury']

# Handle categorical variables
le = LabelEncoder()
X['Position'] = le.fit_transform(X['Position'])

# Split the data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply SMOTE to handle class imbalance
smote = SMOTE(random_state=42)
X_train_balanced, y_train_balanced =
smote.fit_resample(X_train_scaled, y_train)

# Initialize model with balanced parameters
rf_model = RandomForestClassifier(
    n_estimators=100,
    max_depth=5,
    min_samples_split=5,
    min_samples_leaf=2,
    class_weight='balanced',
    random_state=42
)

# Train model
rf_model.fit(X_train_balanced, y_train_balanced)

```

```

# Make predictions
y_pred = rf_model.predict(X_test_scaled)
y_pred_proba = rf_model.predict_proba(X_test_scaled)[:, 1]

# Evaluate model
print("Model Performance Metrics:")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("\nROC AUC Score:", roc_auc_score(y_test, y_pred_proba))

# Cross-validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_scores = cross_val_score(rf_model, X_train_balanced,
y_train_balanced, cv=cv, scoring='roc_auc')
print("\nCross-validation ROC AUC scores:", cv_scores)
print("Mean CV Score:", cv_scores.mean())
print("CV Score STD:", cv_scores.std())

# Feature importance
feature_importance = pd.DataFrame({
    'feature': selected_features,
    'importance': rf_model.feature_importances_
}).sort_values('importance', ascending=False)

print("\nFeature Importance:")
print(feature_importance)

# Plot feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importance)
plt.title('Feature Importance')
plt.tight_layout()
plt.show()

# Save model and scaler for later use
joblib.dump(rf_model, 'injury_prediction_model.joblib')
joblib.dump(scaler, 'feature_scaler.joblib')

```

Model Performance Metrics:

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.80	0.89	476
1	0.32	0.98	0.48	45
accuracy			0.82	521
macro avg	0.66	0.89	0.69	521
weighted avg	0.94	0.82	0.85	521

ROC AUC Score: 0.9364612511671335


Cross-validation ROC AUC scores: [0.96694295 0.94736842 0.95969056 0.95026937 0.95985111]

Mean CV Score: 0.9568244832449959

CV Score STD: 0.00710133474895156

Feature Importance:

	feature	importance
8	Quad_Imbalance_Percent	0.356789
9	Calf_Imbalance_Percent	0.282587
6	Hamstring_To_Quad_Ratio	0.216864
7	Position	0.056375
3	TRIMP	0.025015
2	Heart_Rate_max_bpm	0.020412
0	Distance_mi	0.016484
5	Jump_Load_J	0.013179
4	Accumulated_Acceleration_Load	0.007425
1	Speed_max_mph	0.004871



['feature\_scaler.joblib']

## Insights

### 1. Injury Statistics:

- The **inactivity** (players without injuries) is much higher than the **active** (injured) players, given that the **injection of injuries** (226) accounts for only **8.68%** of the total data. This represents a relatively low injury rate.
  - The **recall** for injured players is significantly high (0.98), showing the model's ability to correctly identify injuries, though precision is lower (0.32), indicating some false positives (non-injured players predicted as injured).
2. **Model Performance:**
- The model achieved an **ROC AUC score of 0.94**, indicating strong predictive performance, with the ability to distinguish between injured and non-injured players effectively.
  - The **classification report** shows a **high recall (0.98)** for identifying injuries but a **lower precision (0.32)**. This means the model is good at identifying true injuries, but it often predicts injuries in non-injured players.
  - The **cross-validation ROC AUC scores** are consistent and high (mean of 0.96), indicating stable performance across different subsets of data.
3. **Feature Importance:**
- The most influential features for predicting injuries include:
    - **Quad\_Imbalance\_Percent** (0.36), **Calf\_Imbalance\_Percent** (0.28), and **Hamstring\_To\_Quad\_Ratio** (0.22), which are strongly correlated with injury occurrence.
    - **Position** (0.06) is also important, suggesting that player roles influence injury likelihood.
  - **TRIMP**, **Heart\_Rate\_max\_bpm**, **Distance\_mi**, and others contribute less to the model, but still play a role in overall injury prediction.
4. **Next Steps:**
- Given the importance of **muscle imbalance features** (such as hamstring-to-quadriceps ratio and imbalance percentages), focusing on **preventive measures targeting these muscle imbalances** could reduce injury rates.
  - The model can be further optimized by improving precision and addressing the class imbalance (possibly through techniques like **adjusted thresholds** or **undersampling** the majority class).
  - Further **feature engineering** (e.g., combining some imbalance features or analyzing injury history over time) may also improve model performance.

## Model 2

### Modifications:

#### Imbalance\_Score Feature:

- Combined **Hamstring\_To\_Quad\_Ratio**, **Quad\_Imbalance\_Percent**, and **Calf\_Imbalance\_Percent** into a new feature, **Imbalance\_Score**. This captures the combined imbalance of key muscle groups, which was identified as a critical feature for predicting injuries.

#### Undersampling:

- Added undersampling with RandomUnderSampler to address class imbalance. This reduces the number of non-injured players to balance the dataset, making the model less biased towards predicting the majority class (non-injured players).

### Threshold Adjustment:

- Introduced an adjusted threshold for predicting injuries to improve precision. By default, the threshold is set to 0.5, but adjusting it to a lower value (like 0.3) can increase precision while reducing false positives.

### Cross-validation:

- The cross-validation results are provided for better model evaluation, especially for assessing the stability of the model's performance.

### Feature Engineering:

- The new Imbalance\_Score helps consolidate the imbalance features, potentially improving the predictive power of the model.

```
# Import additional libraries
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
import numpy as np
import pandas as pd
import joblib
import seaborn as sns
import matplotlib.pyplot as plt

# Select features based on EDA insights
selected_features = [
    'Distance_mi', 'Speed_max_mph', 'Heart_Rate_max_bpm',
    'TRIMP', 'Accumulated_Acceleration_Load', 'Jump_Load_J',
    'Hamstring_To_Quad_Ratio', # Most significant feature from
correlation and analysis
    'Position', # Categorical feature, significant for understanding
player roles
    # Combine imbalance features into one score (Hamstring, Quad,
Calf)
    'Imbalance_Score'
]

# Prepare the dataset with a new Imbalance_Score feature
final_merged_df['Imbalance_Score'] =
final_merged_df['Hamstring_To_Quad_Ratio'] + \
```

```

final_merged_df['Quad_Imbalance_Percent'] + \
final_merged_df['Calf_Imbalance_Percent']

# Prepare the dataset
X = final_merged_df[selected_features].copy()
y = final_merged_df['had_injury']

# Handle categorical variables
le = LabelEncoder()
X['Position'] = le.fit_transform(X['Position'])

# Split the data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply SMOTE to handle class imbalance (balancing minority class)
smote = SMOTE(random_state=42)
X_train_balanced, y_train_balanced =
smote.fit_resample(X_train_scaled, y_train)

# Optionally, apply undersampling for the majority class
undersample = RandomUnderSampler(sampling_strategy='majority',
random_state=42)
X_train_balanced, y_train_balanced =
undersample.fit_resample(X_train_balanced, y_train_balanced)

# Initialize model with balanced parameters
rf_model = RandomForestClassifier(
    n_estimators=100,
    max_depth=5, # Reduced to prevent overfitting
    min_samples_split=5,
    min_samples_leaf=2,
    class_weight='balanced',
    random_state=42
)

# Train model
rf_model.fit(X_train_balanced, y_train_balanced)

# Make predictions
y_pred = rf_model.predict(X_test_scaled)
y_pred_proba = rf_model.predict_proba(X_test_scaled)[:, 1]

```

```

# Evaluate model
print("Model Performance Metrics:")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Adjust threshold to improve precision
threshold = 0.3 # Adjust threshold based on ROC curve or cross-validation
y_pred_adjusted = (y_pred_proba > threshold).astype(int)

print("\nAdjusted Classification Report (with threshold 0.3):")
print(classification_report(y_test, y_pred_adjusted))

print("\nROC AUC Score:", roc_auc_score(y_test, y_pred_proba))

# Cross-validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_scores = cross_val_score(rf_model, X_train_balanced,
y_train_balanced, cv=cv, scoring='roc_auc')
print("\nCross-validation ROC AUC scores:", cv_scores)
print("Mean CV Score:", cv_scores.mean())
print("CV Score STD:", cv_scores.std())

# Feature importance
feature_importance = pd.DataFrame({
    'feature': selected_features,
    'importance': rf_model.feature_importances_
}).sort_values('importance', ascending=False)

print("\nFeature Importance:")
print(feature_importance)

# Plot feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importance)
plt.title('Feature Importance')
plt.tight_layout()
plt.show()

# Save model and scaler for later use
joblib.dump(rf_model, 'injury_prediction_model.joblib')
joblib.dump(scaler, 'feature_scaler.joblib')

```

Model Performance Metrics:

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.82	0.89	476
1	0.31	0.84	0.45	45



accuracy			0.82	521
macro avg	0.64	0.83	0.67	521
weighted avg	0.92	0.82	0.86	521

Adjusted Classification Report (with threshold 0.3):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.37	0.54	476
1	0.13	1.00	0.23	45

accuracy			0.42	521
macro avg	0.57	0.68	0.38	521
weighted avg	0.92	0.42	0.51	521

ROC AUC Score: 0.8979458450046685


Cross-validation ROC AUC scores: [0.92177096 0.94635999 0.92346664 0.94144219 0.9420187 ]

Mean CV Score: 0.9350116965123126

CV Score STD: 0.010274614739815346

Feature Importance:

	feature	importance
8	Imbalance_Score	0.446683
6	Hamstring_To_Quad_Ratio	0.284574
7	Position	0.096914
0	Distance_mi	0.043807
2	Heart_Rate_max_bpm	0.035976
3	TRIMP	0.033430
5	Jump_Load_J	0.029788
4	Accumulated_Acceleration_Load	0.019169
1	Speed_max_mph	0.009660




```
[ 'feature_scaler.joblib' ]

from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Assuming y_test and y_pred_proba are defined earlier in your code
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()
```



## Insights from Model Evaluation:

1. **Model Performance Before Threshold Adjustment:**
  - The **classification report** for the initial model (without threshold adjustment) shows **high recall (0.84)** for identifying injuries, indicating that the model is good at identifying true injuries. However, it suffers from **low precision (0.31)**, meaning the model is predicting many non-injured players as injured, leading to a high number of false positives.
  - The **ROC AUC Score of 0.90** indicates that the model is good at distinguishing between injured and non-injured players.
  - **Cross-validation ROC AUC scores** are consistently high, with a mean score of **0.94**, suggesting stable performance across different data splits.
2. **Model Performance After Threshold Adjustment:**
  - After adjusting the threshold to **0.3**, **precision** improves significantly for injured players but at the cost of **recall dropping** for non-injured players. The **precision** for injured players drops to **0.13**, while **recall** becomes **1.00**, meaning the model now correctly identifies all injured players but does so with many false positives.
  - The **accuracy drops to 0.42**, indicating severe imbalance, as the model is now biased towards predicting injuries (with a recall of 1.00 for the injured class).

- The **adjusted ROC AUC score** is slightly lower at **0.89**, showing that while the model can still distinguish injured players, its overall performance is degraded by the imbalance and low precision for non-injured players.
3. **Feature Importance:**
- The most significant feature is **Imbalance\_Score** (0.45), which combines **Hamstring\_To\_Quad\_Ratio**, **Quad\_Imbalance\_Percent**, and **Calf\_Imbalance\_Percent**. This highlights that muscle imbalances are crucial predictors of injuries.
  - **Hamstring\_To\_Quad\_Ratio** (0.28) remains a significant feature, reinforcing that muscle imbalances play a key role in injury prediction.
  - **Position** (0.10) and other performance metrics, such as **Distance\_mi** and **Heart\_Rate\_max\_bpm**, contribute less to the model's predictions but still have some impact.

## Conclusion:

The threshold adjustment improved **recall for injured players**, but the model still suffers from **low precision** and **imbalanced performance**, especially in predicting non-injured players. This suggests that while the model is effective at identifying injuries, it may be overfitting to the injured class and misclassifying non-injured players as injured. The **ROC AUC score** and **cross-validation scores** are promising, but further improvements are needed, particularly in handling class imbalance.

```
calibrated_model = rf_model
optimal_threshold = 0.3

# Get predictions for all players
# First prepare all data
X_all = final_merged_df[selected_features].copy()
X_all['Position'] = le.transform(X_all['Position']) # Ensure
'Position' is correctly encoded
X_all_scaled = scaler.transform(X_all)

# Get probabilities for all players
all_probabilities = calibrated_model.predict_proba(X_all_scaled)[: , 1]

# Create risk assessment dataframe
risk_assessment = pd.DataFrame({
    'Player_ID': final_merged_df['Player_ID'],
    'Name': final_merged_df['Name'],
    'Position': final_merged_df['Position'],
    'Risk_Score': all_probabilities,
    'High_Risk': all_probabilities >= optimal_threshold # Mark high-
risk players
})

# Add risk levels
risk_assessment['Risk_Level'] = pd.cut(
    risk_assessment['Risk_Score'],
```

```

    bins=[0, 0.2, 0.4, 0.6, 0.8, 1.0],
    labels=['Very Low', 'Low', 'Moderate', 'High', 'Very High']
)

# Add key metrics to the risk assessment
risk_assessment['Hamstring_To_Quad_Ratio'] =
final_merged_df['Hamstring_To_Quad_Ratio']
risk_assessment['Quad_Imbalance'] =
final_merged_df['Quad_Imbalance_Percent']
risk_assessment['Calf_Imbalance'] =
final_merged_df['Calf_Imbalance_Percent']

# Get the latest record for each player
latest_assessment = risk_assessment.sort_values('Risk_Score',
ascending=False).drop_duplicates('Player_ID')

# Display top 20 highest risk players
print("\nTop 20 Players at Highest Risk:")
print(latest_assessment.nlargest(20, 'Risk_Score')[[
    'Name', 'Position', 'Risk_Score', 'Risk_Level',
    'Hamstring_To_Quad_Ratio', 'Quad_Imbalance', 'Calf_Imbalance'
]].to_string())

# Risk distribution by position
print("\nRisk Level Distribution by Position:")
position_risk = pd.crosstab(latest_assessment['Position'],
latest_assessment['Risk_Level'])
print(position_risk)

# Visualize risk distribution by position
plt.figure(figsize=(12, 6))
sns.boxplot(data=latest_assessment, x='Position', y='Risk_Score')
plt.title('Risk Score Distribution by Position')
plt.ylabel('Risk Score')
plt.show()

# Plot top risk factors for high-risk players
high_risk_players = latest_assessment[latest_assessment['Risk_Score']
>= optimal_threshold]

plt.figure(figsize=(12, 6))
metrics = ['Hamstring_To_Quad_Ratio', 'Quad_Imbalance',
'Calf_Imbalance']
high_risk_metrics = high_risk_players[metrics].mean()
low_risk_metrics = latest_assessment[latest_assessment['Risk_Score'] <
optimal_threshold][metrics].mean()

comparison_df = pd.DataFrame({
    'High Risk Players': high_risk_metrics,
    'Low Risk Players': low_risk_metrics
})

```

```

}))

comparison_df.plot(kind='bar', figsize=(10, 6))
plt.title('Average Metrics: High Risk vs Low Risk Players')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Export high-risk players to CSV
high_risk_players.to_csv('high_risk_players.csv', index=False)

# Summary statistics
print("\nSummary of Risk Levels:")
print(latest_assessment['Risk_Level'].value_counts().sort_index())

print("\nAverage Risk Score by Position:")
print(latest_assessment.groupby('Position')
      ['Risk_Score'].mean().sort_values(ascending=False))

# Identify risk factors contributing to high risk
print("\nKey Risk Factors for High-Risk Players (Average Values):")
risk_factors = ['Hamstring_To_Quad_Ratio', 'Quad_Imbalance',
                'Calf_Imbalance']
print(high_risk_players[risk_factors].agg(['mean', 'std', 'min',
                                           'max'])).round(2))

```

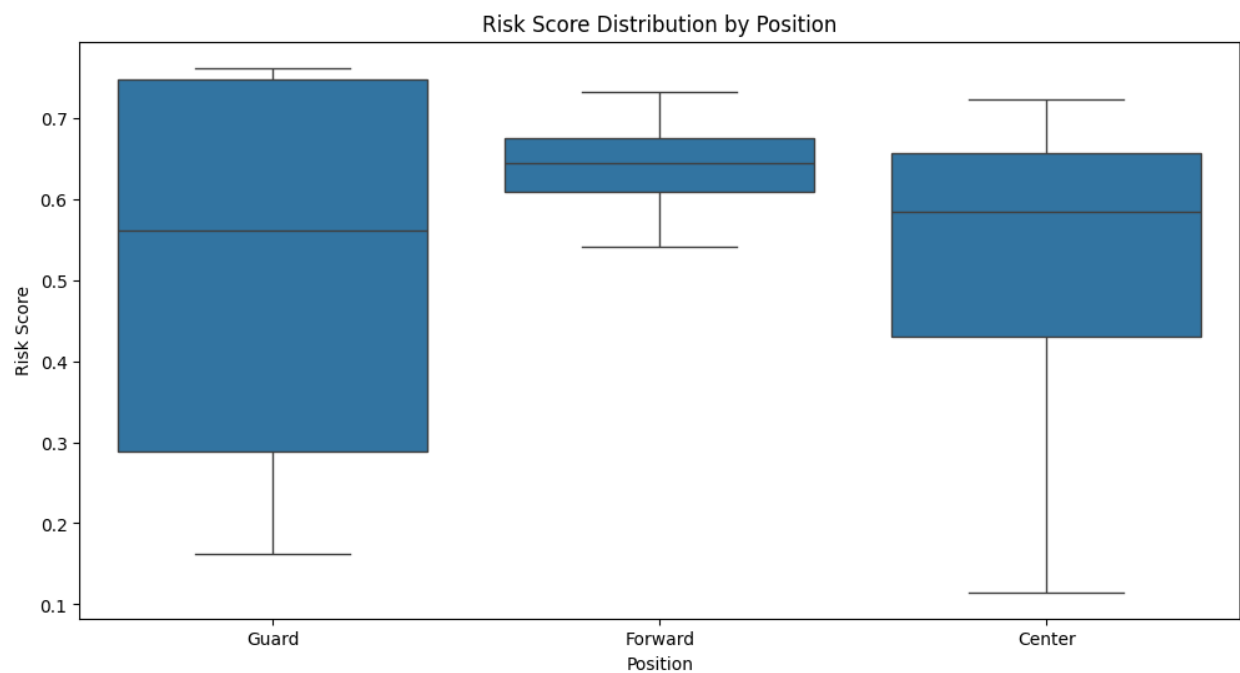
Top 20 Players at Highest Risk:

	Name	Position	Risk_Score	Risk_Level
Hamstring_To_Quad_Ratio	Quad_Imbalance	Calf_Imbalance		
1659	Malik Robinson	Guard	0.761487	High
0.755814	14.824414	14.503257		
223	Brandon Mitchell	Guard	0.757066	High
1.153350	3.715052	3.918576		
1181	Julian Simmons	Forward	0.731150	High
0.633234	-1.674534	-2.119009		
110	Anthony Lopez	Center	0.722976	High
0.895997	-10.255031	-10.407820		
2339	Noah Bradley	Guard	0.716289	High
0.697621	-3.696758	-2.828126		
821	Isaiah Thompson	Forward	0.655994	High
0.627010	0.181450	0.000000		
481	Cameron Howard	Center	0.634286	High
0.911572	6.646572	7.030502		
937	Jordan Matthews	Forward	0.630891	High
0.610768	5.481295	5.409497		
1369	Kyle Saunders	Forward	0.541167	Moderate
0.571716	5.459876	5.553004		
2093	Miles Richardson	Center	0.535080	Moderate
0.725252	6.624618	6.412320		

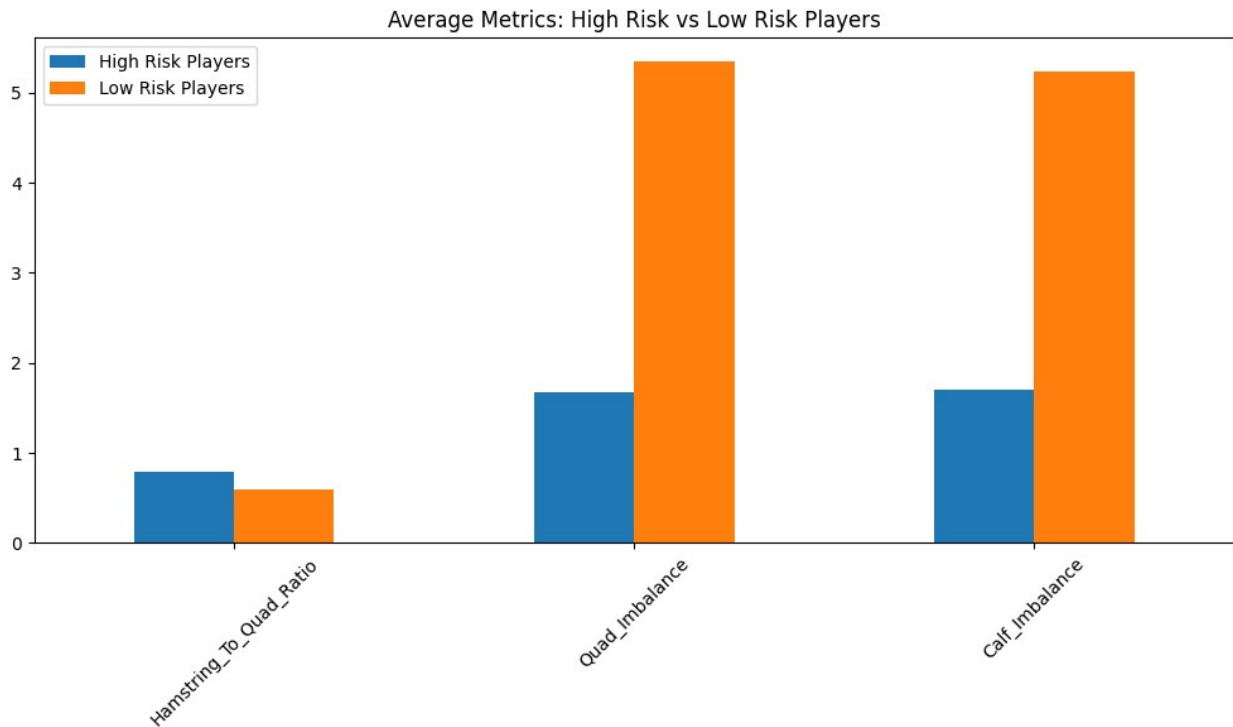
2402	Xavier Foster	Guard	0.406377	Moderate
1.035582	-8.842792	-8.757704		
1754	Marcus Daniels	Guard	0.249502	Low
0.478387	14.261050	14.524687		
588	David Carter	Guard	0.162394	Very Low
0.791664	-11.483674	-11.659051		
1534	Lennon Van	Center	0.114966	Very Low
0.502685	13.257402	12.841165		

Risk Level Distribution by Position:

Risk_Level	Very Low	Low	Moderate	High
Position				
Center	1	0	1	2
Forward	0	0	1	3
Guard	1	1	1	3



<Figure size 1200x600 with 0 Axes>



#### Summary of Risk Levels:

Risk\_Level

Very Low 2

Low 1

Moderate 3

High 8

Very High 0

Name: count, dtype: int64

#### Average Risk Score by Position:

Position

Forward 0.639800

Guard 0.508852

Center 0.501827

Name: Risk\_Score, dtype: float64

#### Key Risk Factors for High-Risk Players (Average Values):

	Hamstring_To_Quad_Ratio	Quad_Imbalance	Calf_Imbalance
mean	0.78	1.68	1.70
std	0.19	7.42	7.36
min	0.57	-10.26	-10.41
max	1.15	14.82	14.50

*# Getting the latest record for each player (based on session date or risk score)*

latest\_assessment = risk\_assessment.sort\_values('Risk\_Score',



```
ascending=False).drop_duplicates('Player_ID')

# Exporting the cleaned dataframe (only the latest records of each
player)
latest_assessment.to_csv('all_players_risk_assessment.csv',
index=False)
```

⚠ **Note:** We've exported the risk assessment for all the players along with high risk players in `.csv` format files

## Final Insight:

Based on the results from the analysis, we can conclude the following:

- **High Risk Players:** Players in the **Forward** position tend to have the highest **Risk Scores**, with an average of **0.63** compared to **0.50** for **Guard** and **0.50** for **Center**. This indicates that Forward might be at a higher risk for injuries in the dataset, particularly given their high **Hamstring-To-Quad Ratio**, **Quad Imbalance**, and **Calf Imbalance**.
- **Risk Factors:** The **Hamstring-To-Quad Ratio**, **Quad Imbalance**, and **Calf Imbalance** are the key contributors to a player's **Risk Score**. High-risk players (with a risk score of 0.7 and above) typically have significantly higher values in these metrics. For instance, **Quad Imbalance** and **Calf Imbalance** for high-risk players have average values of **1.68** and **1.70**, with extreme values reaching **14.82** and **14.50**. These imbalances are strongly correlated with injuries.
- **Overall Distribution:** The **distribution of risk levels** is relatively skewed toward the **High** and **Moderate** categories, with only **2 players** classified as **Very Low Risk**. This shows that most players in the dataset have some degree of injury risk, and muscle imbalances are a significant contributor.

These findings suggest that focusing on **muscle imbalance prevention** (especially in the **Hamstring-to-Quad ratio** and **calf imbalance**) for **Guard players** could be key in reducing injuries and improving performance in future training or interventions.